

# Database Design Document

## Job Portal with Smart Resume–Job Matching

### 1. Overview

This document details the database design for the Job Portal application. The database selected is **PostgreSQL**. It stores users (candidates & recruiters), resumes, job postings, applications, AI matching results, and authentication tokens. The design focuses on data integrity, security, and performance for a web-scale mini-project.

### 2. Database Technology

- Primary DBMS: PostgreSQL (12+ recommended)
- Access: via SQLAlchemy ORM (FastAPI integration) or async ORM (Tortoise/SQLModel)
- Backups: Regular pg\_dump exports and point-in-time recovery if using WAL archiving.

### 3. Entity-Relationship (ER) Overview

Main entities:

- users (candidate or recruiter)
- resumes (one per candidate, or multiple revisions)
- jobs (posted by recruiters)
- applications (candidate applies to a job)
- recommendations (AI similarity results stored for fast retrieval)
- tokens (refresh tokens / sessions)

Relationships:

- users (1) — (M) resumes
- users (recruiter) (1) — (M) jobs
- users (candidate) (1) — (M) applications — (M) jobs (many-to-many via applications)
- jobs (1) — (M) recommendations (each job can have many recommended candidates)

### 4. Table Schemas

Below are recommended table definitions with columns, types, constraints, and indexes. Use UUIDs for primary keys in production if preferred.

#### 4.1 users

```
CREATE TABLE users ( id SERIAL PRIMARY KEY, role VARCHAR(20) NOT NULL CHECK (role IN ('candidate','recruiter','admin')), name VARCHAR(200) NOT NULL, email VARCHAR(255) UNIQUE NOT NULL, password_hash VARCHAR(255) NOT NULL, profile_json JSONB, -- flexible profile fields is_active BOOLEAN DEFAULT TRUE, created_at TIMESTAMP WITH TIME ZONE DEFAULT now(), updated_at TIMESTAMP WITH TIME ZONE DEFAULT now() ); -- Indexes CREATE INDEX idx_users_email ON users(email);
```

#### 4.2 resumes

```
CREATE TABLE resumes ( id SERIAL PRIMARY KEY, user_id INTEGER REFERENCES users(id) ON DELETE CASCADE, filename VARCHAR(255), storage_path VARCHAR(512), -- S3/Server path text_content TEXT, -- extracted resume text parsed_skills TEXT[], -- optional array of skills uploaded_at TIMESTAMP WITH TIME ZONE DEFAULT now(), is_active BOOLEAN DEFAULT TRUE ); CREATE INDEX idx_resumes_user ON resumes(user_id); CREATE INDEX idx_resumes_skills ON resumes USING GIN(parsed_skills);
```

#### 4.3 jobs

```
CREATE TABLE jobs ( id SERIAL PRIMARY KEY, recruiter_id INTEGER REFERENCES users(id) ON DELETE SET NULL, title VARCHAR(255) NOT NULL, description TEXT, required_skills TEXT[], -- array of skills location VARCHAR(255), employment_type VARCHAR(50), created_at TIMESTAMP WITH TIME ZONE DEFAULT now(), updated_at TIMESTAMP WITH TIME ZONE DEFAULT now(), is_active BOOLEAN DEFAULT TRUE );  
CREATE INDEX idx_jobs_recruiter ON jobs(recruiter_id); CREATE INDEX idx_jobs_skills ON jobs USING GIN(required_skills);
```

## 4.4 applications

```
CREATE TABLE applications ( id SERIAL PRIMARY KEY, job_id INTEGER REFERENCES jobs(id) ON DELETE CASCADE, candidate_id INTEGER REFERENCES users(id) ON DELETE CASCADE, resume_id INTEGER REFERENCES resumes(id) ON DELETE SET NULL, cover_letter TEXT, status VARCHAR(50) DEFAULT 'applied', -- applied, shortlisted, rejected, interviewed, hired applied_at TIMESTAMP WITH TIME ZONE DEFAULT now() );  
CREATE INDEX idx_applications_job ON applications(job_id);  
CREATE INDEX idx_applications_candidate ON applications(candidate_id);
```

## 4.5 recommendations

```
CREATE TABLE recommendations ( id SERIAL PRIMARY KEY, job_id INTEGER REFERENCES jobs(id) ON DELETE CASCADE, candidate_id INTEGER REFERENCES users(id) ON DELETE CASCADE, resume_id INTEGER REFERENCES resumes(id), similarity_score REAL NOT NULL, computed_at TIMESTAMP WITH TIME ZONE DEFAULT now() );  
CREATE INDEX idx_recommendations_job ON recommendations(job_id);  
CREATE INDEX idx_recommendations_candidate ON recommendations(candidate_id);  
CREATE INDEX idx_recommendations_score ON recommendations(job_id, similarity_score DESC);
```

## 4.6 tokens

```
CREATE TABLE tokens ( id SERIAL PRIMARY KEY, user_id INTEGER REFERENCES users(id) ON DELETE CASCADE, refresh_token_hash VARCHAR(255) NOT NULL, user_agent TEXT, ip_address INET, created_at TIMESTAMP WITH TIME ZONE DEFAULT now(), expires_at TIMESTAMP WITH TIME ZONE );  
CREATE INDEX idx_tokens_user ON tokens(user_id);
```

## 5. Relationships & Constraints

- Foreign keys enforce referential integrity (ON DELETE behaviours as suggested).
- Use CHECK constraints for enums (roles, application status).
- Use JSONB for flexible profile fields and future-proofing.
- Use GIN indexes for array/JSONB fields to speed up skill-based queries.

## 6. Security & Optimization

Security best practices:

- Use parameterized queries / ORM to prevent SQL injection.
- Principle of least privilege: database user has minimal rights (no superuser for app).
- Encrypt sensitive data at rest and in transit (TLS).
- Hash refresh tokens and passwords (bcrypt) — never store plaintext.
- Validate and sanitize all inputs in backend.

Optimization:

- Add indexes on frequently queried columns (email, job\_id, candidate\_id).
- Use GIN indexes for full-text search or array fields (skills).
- Consider materialized views for expensive aggregations (e.g., top jobs).
- Partition large tables if data grows (applications, recommendations).
- Use EXPLAIN ANALYZE to optimize slow queries.

## 7. Backup & Maintenance

- Regular full backups: pg\_dumpall or pg\_basebackup.
- Point-in-time recovery (PITR) via WAL archiving for critical systems.
- Routine vacuum & analyze to maintain stats.
- Monitor disk, slow queries, connections; set alerts.
- Test backups periodically by restoring to staging environment.

## 8. Example Queries

```
-- 1. Fetch top 10 recommended candidates for a job (from recommendations table)
SELECT r.candidate_id,
u.name, r.similarity_score
FROM recommendations r
JOIN users u ON u.id = r.candidate_id
WHERE r.job_id = 10
ORDER BY r.similarity_score DESC LIMIT 10;
-- 2. Get recommended jobs for a candidate (join
recommendations and jobs)
SELECT j.id, j.title, rec.similarity_score
FROM recommendations rec
JOIN jobs j ON j.id = rec.job_id
WHERE rec.candidate_id = 5
ORDER BY rec.similarity_score DESC;
-- 3. Full-text search on job
descriptions (requires tsvector column/index)
SELECT id, title
FROM jobs
WHERE to_tsvector('english',
description) @@ plainto_tsquery('python developer');
```

## 9. FastAPI Integration (SQLAlchemy models)

Use SQLAlchemy (or SQLAlchemy) to define ORM models. Example SQLAlchemy model for Job and Recommendation simplified below:

```
from sqlalchemy import Column, Integer, String, Text, ARRAY, ForeignKey, Float, DateTime, func
from sqlalchemy.orm import relationship
from database import Base
class Job(Base):
    __tablename__ = 'jobs'
    id = Column(Integer, primary_key=True, index=True)
    recruiter_id = Column(Integer, ForeignKey('users.id'))
    title = Column(String, nullable=False)
    description = Column(Text)
    required_skills = Column(ARRAY(String))
    created_at = Column(DateTime(timezone=True), server_default=func.now())
    recommendations = relationship('Recommendation', back_populates='job')
class Recommendation(Base):
    __tablename__ = 'recommendations'
    id = Column(Integer, primary_key=True, index=True)
    job_id = Column(Integer, ForeignKey('jobs.id'))
    candidate_id = Column(Integer, ForeignKey('users.id'))
    resume_id = Column(Integer)
    similarity_score = Column(Float)
    computed_at = Column(DateTime(timezone=True), server_default=func.now())
```

## 10. Closing Notes

This design aims to be practical for a 3rd-year mini-project while following best practices for security, performance, and maintainability. You can extend the schema with additional tables for interviews, company profiles, analytics, and audit logs as needed.